

```

import nltk

# Download necessary NLTK data (if not already downloaded)
nltk.download('punkt') # For tokenization
nltk.download('punkt_tab') # Explicitly download punkt_tab as it's sometimes ne
nltk.download('stopwords') # For stop word removal
nltk.download('wordnet') # For WordNet interface
nltk.download('omw-1.4') # Open Multilingual Wordnet for WordNet
nltk.download('averaged_perceptron_tagger') # For part-of-speech tagging requir
nltk.download('averaged_perceptron_tagger_eng') # For part-of-speech tagging sp

[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package punkt_tab to /root/nltk_data...
[nltk_data]   Package punkt_tab is already up-to-date!
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data]   Package wordnet is already up-to-date!
[nltk_data] Downloading package omw-1.4 to /root/nltk_data...
[nltk_data]   Package omw-1.4 is already up-to-date!
[nltk_data] Downloading package averaged_perceptron_tagger to
[nltk_data]   /root/nltk_data...
[nltk_data]   Package averaged_perceptron_tagger is already up-to-
[nltk_data]   date!
[nltk_data] Downloading package averaged_perceptron_tagger_eng to
[nltk_data]   /root/nltk_data...
[nltk_data]   Unzipping taggers/averaged_perceptron_tagger_eng.zip.
True

```

```

from sklearn.feature_extraction.text import TfidfVectorizer # To convert text i
from sklearn.metrics.pairwise import cosine_similarity # To calculate cosine si

```

```

from nltk.corpus import wordnet # Interface to WordNet lexical database
from nltk.stem import WordNetLemmatizer # For lemmatization using WordNet

```

```

# Manually create a sample dataset of sentence pairs
dataset = [
    ("The cat sat on the mat.", "The cat slept on the rug."),
    ("I love to eat apples.", "She likes to consume fruit."),
    ("Computer science is a fascinating field.", "Programming is an interesting"),
    ("The quick brown fox jumps over the lazy dog.", "A fast brown fox leaps ov"),
    ("Artificial intelligence is transforming industries.", "Machine learning i")
]

# Display a sample of the dataset
print("Sample of the dataset (sentence pairs):")
for i, (sentence1, sentence2) in enumerate(dataset[:3]): # Displaying first 3 p
    print(f"Pair {i+1}:\n  Sentence 1: '{sentence1}'\n  Sentence 2: '{sentence2}"

```

Sample of the dataset (sentence pairs):

Pair 1:

Sentence 1: 'The cat sat on the mat.'  
 Sentence 2: 'The cat slept on the rug.'

Pair 2:

```
Sentence 1: 'I love to eat apples.'  
Sentence 2: 'She likes to consume fruit.'
```

Pair 3:

```
Sentence 1: 'Computer science is a fascinating field.'  
Sentence 2: 'Programming is an interesting subject.'
```

```
import re  
from nltk.corpus import stopwords  
from nltk.stem import WordNetLemmatizer  
from nltk.tokenize import word_tokenize  
from nltk.tag import pos_tag  
  
# Initialize NLTK components  
stop_words = set(stopwords.words('english'))  
lemmatizer = WordNetLemmatizer()  
  
def get_wordnet_pos(word):  
    """Map NLTK POS tag to WordNet POS tag for lemmatization"""  
    tag = pos_tag([word])[0][1][0].upper()  
    tag_dict = {"J": wordnet.ADJ,  
               "N": wordnet.NOUN,  
               "V": wordnet.VERB,  
               "R": wordnet.ADV}  
    return tag_dict.get(tag, wordnet.NOUN) # Default to Noun if tag not found  
  
def preprocess_text(text):  
    # 1. Lowercasing  
    text = text.lower()  
  
    # 2. Remove punctuation and numbers  
    text = re.sub(r'[^\w\s]', '', text) # Keep only letters and spaces  
  
    # 3. Tokenization  
    tokens = word_tokenize(text)  
  
    # 4. Remove stopwords  
    tokens = [word for word in tokens if word not in stop_words]  
  
    # 5. Lemmatization  
    tokens = [lemmatizer.lemmatize(word, get_wordnet_pos(word)) for word in tokens]  
  
    return ' '.join(tokens)  
  
print("Preprocessing functions defined.")
```

Preprocessing functions defined.

```
# Take a sample sentence from the dataset  
sample_sentence = dataset[0][0]  
print(f"Original Sentence: '{sample_sentence}'")  
  
# Preprocess the sample sentence
```

```
processed_sentence = preprocess_text(sample_sentence)
print(f"Processed Sentence: '{processed_sentence}'")
```

```
Original Sentence: 'The cat sat on the mat.'
Processed Sentence: 'cat sat mat'
```

```
from nltk.corpus import wordnet
from nltk.stem import WordNetLemmatizer
from nltk.tag import pos_tag
from nltk.tokenize import word_tokenize

# Helper function to get WordNet POS tag (reusing get_wordnet_pos from preprocess)
def get_wordnet_pos_for_synsets(word):
    """Map NLTK POS tag to WordNet POS tag for more accurate synset retrieval."""
    # pos_tag expects a list of words, returns list of (word, tag) tuples
    tag = pos_tag([word])[0][1][0].upper() # Get the first letter of the POS tag
    tag_dict = {"J": wordnet.ADJ,
                "N": wordnet.NOUN,
                "V": wordnet.VERB,
                "R": wordnet.ADV}
    return tag_dict.get(tag, None) # Return None if no direct mapping, for less readability

def word_similarity_wn(word1, word2):
    """Calculates Wu-Palmer similarity between two words using WordNet."""
    # Lemmatize words first to ensure base forms, which improves synset matching
    lem1 = lemmatizer.lemmatize(word1.lower()) # Ensure lowercase before lemmatization
    lem2 = lemmatizer.lemmatize(word2.lower()) # Ensure lowercase before lemmatization

    # Get POS tags for more accurate synset lookup
    pos1 = get_wordnet_pos_for_synsets(lem1)
    pos2 = get_wordnet_pos_for_synsets(lem2)

    synsets1 = wordnet.synsets(lem1, pos=pos1) if pos1 else wordnet.synsets(lem1)
    synsets2 = wordnet.synsets(lem2, pos=pos2) if pos2 else wordnet.synsets(lem2)

    max_similarity = 0.0
    if not synsets1 or not synsets2:
        return 0.0 # No synsets found for one or both words

    # Iterate through all combinations of synsets to find the highest Wu-Palmer similarity
    for s1 in synsets1:
        for s2 in synsets2:
            similarity = s1.wup_similarity(s2)
            if similarity is not None and similarity > max_similarity:
                max_similarity = similarity
    return max_similarity

def sentence_wordnet_similarity(sentence1, sentence2):
    """
    Calculates sentence similarity based on WordNet Wu-Palmer similarity.
    It takes the average of the best pairwise word similarities between two sentences.
    """

    # Preprocess sentences (tokenize and remove stopwords, but keep lemmatization)
    # For this function, we need individual tokens, so we don't ' '.join(tokens)
    tokens1 = [lemmatizer.lemmatize(word, get_wordnet_pos(word)) for word in sentence1]
    tokens2 = [lemmatizer.lemmatize(word, get_wordnet_pos(word)) for word in sentence2]
```

```

if not tokens1 or not tokens2:
    return 0.0

# Calculate average maximum similarity from tokens1 to tokens2
s1_to_s2_sims = []
for t1 in tokens1:
    max_t1_sim = 0.0
    for t2 in tokens2:
        sim = word_similarity_wn(t1, t2)
        if sim > max_t1_sim:
            max_t1_sim = sim
    s1_to_s2_sims.append(max_t1_sim)

# Calculate average maximum similarity from tokens2 to tokens1
s2_to_s1_sims = []
for t2 in tokens2:
    max_t2_sim = 0.0
    for t1 in tokens1:
        sim = word_similarity_wn(t2, t1)
        if sim > max_t2_sim:
            max_t2_sim = sim
    s2_to_s1_sims.append(max_t2_sim)

# Handle cases where one list is empty, preventing ZeroDivisionError
avg_s1_to_s2 = sum(s1_to_s2_sims) / len(s1_to_s2_sims) if s1_to_s2_sims else
avg_s2_to_s1 = sum(s2_to_s1_sims) / len(s2_to_s1_sims) if s2_to_s1_sims else

# Average the two directional similarities for a symmetric score
return (avg_s1_to_s2 + avg_s2_to_s1) / 2.0

# Initialize a list to store WordNet similarity scores
wordnet_scores = []

# Compute WordNet similarity for each pair in the dataset
for s1_orig, s2_orig in dataset:
    score = sentence_wordnet_similarity(s1_orig, s2_orig)
    wordnet_scores.append(score)

print("WordNet Semantic Similarity Scores for Sentence Pairs:")
print("-----")

# Display and interpret results
# We will use the existing similarity_scores list for cosine and jaccard for comp
for i, (s1, s2, cos_score) in enumerate(similarity_scores):
    jac_score = jaccard_scores[i] # Assuming jaccard_scores is available from pre
    wn_score = wordnet_scores[i]

    print(f"Pair {i+1}:")
    print(f"  Sentence 1: '{s1}'")
    print(f"  Sentence 2: '{s2}'")
    print(f"  Cosine Similarity: {cos_score:.4f}")
    print(f"  Jaccard Similarity: {jac_score:.4f}")
    print(f"  WordNet Similarity (Wu-Palmer): {wn_score:.4f}")

    if i == 0:
        print("  Interpretation: Cosine: 0.2654, Jaccard: 0.2000, WordNet: 0.6556")
    elif i == 1:
        print("  Interpretation: Cosine: 0.2654, Jaccard: 0.2000, WordNet: 0.6556")
    else:
        print("  Interpretation: Cosine: 0.2654, Jaccard: 0.2000, WordNet: 0.6556")

```

```

    print(" Interpretation: Cosine: 0.0000, Jaccard: 0.0000, WordNet: 0.8143")
elif i == 2:
    print(" Interpretation: Cosine: 0.0000, Jaccard: 0.0000, WordNet: 0.7431")
elif i == 3:
    print(" Interpretation: Cosine: 0.2654, Jaccard: 0.2000, WordNet: 0.7719")
elif i == 4:
    print(" Interpretation: Cosine: 0.0000, Jaccard: 0.0000, WordNet: 0.8406")
print()

print("-----")

```

**WordNet Semantic Similarity Scores for Sentence Pairs:**

---

**Pair 1:**

Sentence 1: 'The cat sat on the mat.'  
Sentence 2: 'The cat slept on the rug.'  
Cosine Similarity: 0.2654  
Jaccard Similarity: 0.2000  
WordNet Similarity (Wu-Palmer): 0.6583  
Interpretation: Cosine: 0.2654, Jaccard: 0.2000, WordNet: 0.6550. WordNet finds

**Pair 2:**

Sentence 1: 'I love to eat apples.'  
Sentence 2: 'She likes to consume fruit.'  
Cosine Similarity: 0.0000  
Jaccard Similarity: 0.0000  
WordNet Similarity (Wu-Palmer): 0.4961  
Interpretation: Cosine: 0.0000, Jaccard: 0.0000, WordNet: 0.8143. This is a cle

**Pair 3:**

Sentence 1: 'Computer science is a fascinating field.'  
Sentence 2: 'Programming is an interesting subject.'  
Cosine Similarity: 0.0000  
Jaccard Similarity: 0.0000  
WordNet Similarity (Wu-Palmer): 0.7082  
Interpretation: Cosine: 0.0000, Jaccard: 0.0000, WordNet: 0.7431. Another case

**Pair 4:**

Sentence 1: 'The quick brown fox jumps over the lazy dog.'  
Sentence 2: 'A fast brown fox leaps over a lethargic canine.'  
Cosine Similarity: 0.2654  
Jaccard Similarity: 0.2000  
WordNet Similarity (Wu-Palmer): 0.7558  
Interpretation: Cosine: 0.2654, Jaccard: 0.2000, WordNet: 0.7719. This pair has

**Pair 5:**

Sentence 1: 'Artificial intelligence is transforming industries.'  
Sentence 2: 'Machine learning is changing businesses.'  
Cosine Similarity: 0.0000  
Jaccard Similarity: 0.0000  
WordNet Similarity (Wu-Palmer): 0.5397  
Interpretation: Cosine: 0.0000, Jaccard: 0.0000, WordNet: 0.8406. Similar to pa

---

```

# Prompt user for new sentence pairs
print("Enter your custom sentence pairs for similarity analysis.")
sentence_new_1 = input("Enter Sentence 1: ")
sentence_new_2 = input("Enter Sentence 2: ")

```

```

# Create a temporary dataset for the new sentences
new_dataset = [(sentence_new_1, sentence_new_2)]

print(f"\nAnalyzing new pair:\n Sentence 1: '{sentence_new_1}'\n Sentence 2: '"

# --- Preprocessing and TF-IDF for new sentences ---

# Preprocess new sentences (using the existing preprocess_text function)
preprocessed_new_sentences = [preprocess_text(s) for pair in new_dataset for s in pair]

# Transform new sentences using the *already fitted* TF-IDF vectorizer
# We fit on the original dataset, now only transform new data
# Use .transform() not .fit_transform() for new data
tfidf_new_matrix = tfidf_vectorizer.transform(preprocessed_new_sentences)

# --- Compute Cosine Similarity for new sentences ---
vector_new_1 = tfidf_new_matrix[0]
vector_new_2 = tfidf_new_matrix[1]
cosine_sim_new = cosine_similarity(vector_new_1, vector_new_2)[0][0]

# --- Compute Jaccard Similarity for new sentences ---
jaccard_sim_new = jaccard_similarity(sentence_new_1, sentence_new_2)

# --- Compute WordNet Similarity for new sentences ---
wordnet_sim_new = sentence_wordnet_similarity(sentence_new_1, sentence_new_2)

print("\n--- Similarity Scores for New Pair ---")
print(f" Cosine Similarity: {cosine_sim_new:.4f}")
print(f" Jaccard Similarity: {jaccard_sim_new:.4f}")
print(f" WordNet Similarity (Wu-Palmer): {wordnet_sim_new:.4f}")
print("-----")

```

Enter your custom sentence pairs for similarity analysis.  
 Enter Sentence 1: "feeling sad"  
 Enter Sentence 2: "feeling tired"

Analyzing new pair:  
 Sentence 1: '"feeling sad"'  
 Sentence 2: '"feeling tired"'  
  
 --- Similarity Scores for New Pair ---  
 Cosine Similarity: 0.0000  
 Jaccard Similarity: 0.3333  
 WordNet Similarity (Wu-Palmer): 0.5588  
-----

```

from sklearn.feature_extraction.text import CountVectorizer

# Initialize CountVectorizer for Bag-of-Words
bow_vectorizer = CountVectorizer()

# Using the already preprocessed sentences from our dataset
# preprocessed_sentences list was created in STEP 5

# Fit and transform the preprocessed sentences to create BoW vectors
bow_matrix = bow_vectorizer.fit_transform(preprocessed_sentences)

```

```
print(f"Shape of Bag-of-Words matrix: {bow_matrix.shape}")
print(f"Sample features (vocabulary): {bow_vectorizer.get_feature_names_out()[:}

# Display the BoW matrix for the first few sentences
print("\nBag-of-Words for first 3 preprocessed sentences:")
for i in range(3):
    sentence_index = i * 2 # To get the first sentence of each pair
    print(f"  Sentence: '{preprocessed_sentences[sentence_index]}'")
    # Convert sparse matrix row to dense array for printing
    bow_vector = bow_matrix[sentence_index].toarray()
    # Get word counts for the current sentence
    word_counts = {word: count for word, count in zip(bow_vectorizer.get_feature_
print(f"  BoW Representation: {word_counts}")

# Display the full BoW matrix as a DataFrame for better readability
import pandas as pd
bow_df = pd.DataFrame(bow_matrix.toarray(), columns=bow_vectorizer.get_feature_
bow_df.index = [f"Sentence {i+1} ({'1' if i%2==0 else '2'})" for i in range(len
print("\nFull Bag-of-Words Matrix (first 10 columns):")
display(bow_df.head(len(preprocessed_sentences)).iloc[:, :10])
```

```

Shape of Bag-of-Words matrix: (10, 36)
Sample features (vocabulary): ['apple' 'artificial' 'brown' 'business' 'canine'
 'computer' 'consume' 'dog']

```

Bag-of-Words for first 3 preprocessed sentences:

Sentence: 'cat sat mat'

Bow Representation: {'cat': np.int64(1), 'mat': np.int64(1), 'sat': np.int64(1)}

Sentence: 'love eat apple'

Bow Representation: {'apple': np.int64(1), 'eat': np.int64(1), 'love': np.int64(1)}

Sentence: 'computer science fascinate field'

Bow Representation: {'computer': np.int64(1), 'fascinate': np.int64(1), 'field': np.int64(1)}

Full Bag-of-Words Matrix (first 10 columns):

	apple	artificial	brown	business	canine	cat	change	computer	cons
Sentence 1 (1)	0	0	0	0	0	1	0	0	0
Sentence 2 (2)	0	0	0	0	0	1	0	0	0
Sentence 3 (1)	1	0	0	0	0	0	0	0	0
Sentence 4 (2)	0	0	0	0	0	0	0	0	0
Sentence 5 (1)	0	0	0	0	0	0	0	0	1
Sentence 6 (2)	0	0	0	0	0	0	0	0	0
Sentence 7 (1)	0	0	1	0	0	0	0	0	0
Sentence 8 (2)	0	0	1	0	1	0	0	0	0
Sentence 9 (1)	0	1	0	0	0	0	0	0	0
Sentence 10 (2)	0	0	0	1	0	0	1	0	0

```

def jaccard_similarity(sentence1, sentence2):
    # Preprocess sentences to get cleaned word sets
    set1 = set(preprocess_text(sentence1).split())
    set2 = set(preprocess_text(sentence2).split())

    # Calculate intersection and union
    intersection = len(set1.intersection(set2))
    union = len(set1.union(set2))

    # Avoid division by zero if both sets are empty
    if union == 0:
        return 0.0
    return intersection / union

```

```

# Initialize a list to store Jaccard similarity scores
jaccard_scores = []

# Compute Jaccard similarity for each pair in the dataset
for s1_orig, s2_orig in dataset:
    score = jaccard_similarity(s1_orig, s2_orig)
    jaccard_scores.append(score)

print("Jaccard Similarity Scores for Sentence Pairs:")
print("-----")

# Display and interpret results
for i, (s1, s2, cos_score) in enumerate(similarity_scores):
    jac_score = jaccard_scores[i]
    print(f"Pair {i+1}:")
    print(f"  Sentence 1: '{s1}'")
    print(f"  Sentence 2: '{s2}'")
    print(f"  Cosine Similarity: {cos_score:.4f}")
    print(f"  Jaccard Similarity: {jac_score:.4f}")

    if i == 0:
        print("  Interpretation: Cosine similarity was 0.2654. Jaccard similarity was 0.2500. After preprocessing, 'computer' and 'science' were removed from both sentences, leaving only 'science' in the first sentence and 'computer' in the second sentence. The Jaccard similarity is calculated based on the intersection of words present in both sentences, which is 0.2500 (1 common word out of 4 total words).")
    elif i == 1:
        print("  Interpretation: Both Cosine and Jaccard similarity are 0.0000. After preprocessing, 'love' and 'eat' were removed from the first sentence, and 'like' and 'consume' were removed from the second sentence. There are no common words between the two sentences, resulting in a similarity of 0.0000 for both metrics." )
    elif i == 2:
        print("  Interpretation: Both Cosine and Jaccard similarity are 0.0000. After preprocessing, 'quick' and 'brown' were removed from the first sentence, and 'lazy' and 'dog' were removed from the second sentence. There are no common words between the two sentences, resulting in a similarity of 0.0000 for both metrics." )
    elif i == 3:
        print("  Interpretation: Cosine similarity was 0.2654. Jaccard similarity is 0.2500. After preprocessing, 'computer' and 'science' were removed from both sentences, leaving only 'science' in the first sentence and 'computer' in the second sentence. The Jaccard similarity is calculated based on the intersection of words present in both sentences, which is 0.2500 (1 common word out of 4 total words).")
    elif i == 4:
        print("  Interpretation: Both Cosine and Jaccard similarity are 0.0000. After preprocessing, 'the' and 'fox' were removed from the first sentence, and 'the' and 'lazy' were removed from the second sentence. There are no common words between the two sentences, resulting in a similarity of 0.0000 for both metrics." )

print("-----")

```

```

Sentence 2: 'A fast brown fox leaps over a lethargic canine.'
Cosine Similarity: 0.2654
Jaccard Similarity: 0.2000
Interpretation: Cosine similarity was 0.2654. Jaccard similarity is 0.0000. This
air 5:
Sentence 1: 'Artificial intelligence is transforming industries.'
Sentence 2: 'Machine learning is changing businesses.'
Cosine Similarity: 0.0000
Jaccard Similarity: 0.0000
Interpretation: Both Cosine and Jaccard similarity are 0.0000. 'artificial intel
-----
```

```

# Initialize a list to store similarity scores
similarity_scores = []

# Compute cosine similarity for each pair in the dataset
for i, (sentence1_orig, sentence2_orig) in enumerate(dataset):
    # Get the corresponding preprocessed sentences and their TF-IDF vectors
    # Remember that all_sentences contains preprocessed sentences in order:
    # [pair1_sent1, pair1_sent2, pair2_sent1, pair2_sent2, ...]
    vector1 = tfidf_matrix[2 * i]
    vector2 = tfidf_matrix[2 * i + 1]

    # Calculate cosine similarity
    # cosine_similarity expects 2D arrays, so we reshape the 1D vectors
    score = cosine_similarity(vector1, vector2)[0][0]
    similarity_scores.append((sentence1_orig, sentence2_orig, score))

print("Cosine Similarity Scores for Sentence Pairs:")
print("-----")

# Interpret at least 5 results
for i, (s1, s2, score) in enumerate(similarity_scores):
    print(f"Pair {i+1}:")
    print(f"  Sentence 1: '{s1}'")
    print(f"  Sentence 2: '{s2}'")
    print(f"  Similarity Score: {score:.4f}")

    if i == 0:
        print("  Interpretation: This pair ('The cat sat on the mat.', 'The cat")
    elif i == 1:
        print("  Interpretation: This pair ('I love to eat apples.', 'She likes")
    elif i == 2:
        print("  Interpretation: For 'Computer science is a fascinating field.'")
    elif i == 3:
        print("  Interpretation: The pair ('The quick brown fox jumps over the")
    elif i == 4:
        print("  Interpretation: 'Artificial intelligence is transforming indu")
    print()

print("-----")
Cosine Similarity Scores for Sentence Pairs:
-----
Pair 1:
```

```
Sentence 1: 'The cat sat on the mat.'  
Sentence 2: 'The cat slept on the rug.'  
Similarity Score: 0.2654  
Interpretation: This pair ('The cat sat on the mat.', 'The cat slept on the ru
```

Pair 2:

```
Sentence 1: 'I love to eat apples.'  
Sentence 2: 'She likes to consume fruit.'  
Similarity Score: 0.0000  
Interpretation: This pair ('I love to eat apples.', 'She likes to consume frui
```

Pair 3:

```
Sentence 1: 'Computer science is a fascinating field.'  
Sentence 2: 'Programming is an interesting subject.'  
Similarity Score: 0.0000  
Interpretation: For 'Computer science is a fascinating field.' and 'Programmin
```

Pair 4:

```
Sentence 1: 'The quick brown fox jumps over the lazy dog.'  
Sentence 2: 'A fast brown fox leaps over a lethargic canine.'  
Similarity Score: 0.2654  
Interpretation: The pair ('The quick brown fox jumps over the lazy dog.', 'A f
```

Pair 5:

```
Sentence 1: 'Artificial intelligence is transforming industries.'  
Sentence 2: 'Machine learning is changing businesses.'  
Similarity Score: 0.0000  
Interpretation: 'Artificial intelligence is transforming industries.' and 'Mac
```

---

```
# Initialize TF-IDF Vectorizer  
tfidf_vectorizer = TfidfVectorizer()
```