

```
import os
os.environ['KAGGLE_USERNAME'] = 'sreeyadaripelli'
os.environ['KAGGLE_KEY'] = '5e0752429f497134d5284d136fd300fb'
```

```
import kaggle
kaggle.api.dataset_download_files('crowdflower/twitter-airline-sentiment', path

Dataset URL: https://www.kaggle.com/datasets/crowdflower/twitter-airline-sentiment
```

```
import pandas as pd

tweets_df = pd.read_csv('Tweets.csv')
display(tweets_df.head())

  tweet_id  airline_sentiment  airline_sentiment_confidence  negative
0  570306133677760513           neutral                  1.0000
1  570301130888122368          positive                 0.3486
2  570301083672813571           neutral                 0.6837
3  570301031407624196          negative                 1.0000      Ba...
4  570300817074462722          negative                 1.0000      C...
```

```
print(tweets_df['airline_sentiment'].value_counts())
```

```
airline_sentiment
negative    9178
neutral     3099
positive    2363
Name: count, dtype: int64
```

```
negative_tweets_df = tweets_df[tweets_df['airline_sentiment'] == 'negative']

print(f"Number of negative sentiment tweets: {len(negative_tweets_df)}")
display(negative_tweets_df.head())
```

```
ment tweets: 9178
airline_sentiment  airline_sentiment_confidence  negative
```

	negative	1.0000	B
	negative	1.0000	C
	negative	1.0000	C
	negative	0.6842	L
	negative	1.0000	B

```
from sklearn.feature_extraction.text import TfidfVectorizer

# Assuming 'text_for_tfidf' column exists from a previous preprocessing step
# If it doesn't, this step would need to be preceded by text cleaning and tokenization
vectorizer = TfidfVectorizer(max_features=5000, stop_words='english')
tfidf_matrix = vectorizer.fit_transform(tweets_df['text_for_tfidf'])
feature_names = vectorizer.get_feature_names_out()

print("TF-IDF matrix created with shape:", tfidf_matrix.shape)
print("Number of features (terms):"), len(feature_names)
```

```
TF-IDF matrix created with shape: (14640, 5000)
Number of features (terms):
(None, 5000)
```

```
negative_tweet_indices = negative_tweets_df.index

print(f"Indices of negative tweets (first 5): {negative_tweet_indices[:5].tolist()}")
print(f"Total number of negative tweet indices: {len(negative_tweet_indices)})")

Indices of negative tweets (first 5): [3, 4, 5, 15, 17]
Total number of negative tweet indices: 9178
```

```
negative_tfidf_matrix = tfidf_matrix[negative_tweet_indices]

# Calculate the sum of TF-IDF scores for each term across all negative tweets
sum_tfidf_negative = negative_tfidf_matrix.sum(axis=0)

print(f"Shape of negative TF-IDF matrix: {negative_tfidf_matrix.shape}")
print(f"Shape of sum of TF-IDF scores for negative terms: {sum_tfidf_negative.shape}")
```

```
Shape of negative TF-IDF matrix: (9178, 5000)
Shape of sum of TF-IDF scores for negative terms: (1, 5000)
```

```
import pandas as pd

# Convert the sparse matrix row vector to a dense NumPy array
sum_tfidf_negative_array = sum_tfidf_negative.A.flatten()

# Create a Pandas Series with feature names as index and TF-IDF scores as value
tfidf_scores_negative = pd.Series(sum_tfidf_negative_array, index=feature_names)

# Sort the Series in descending order
tfidf_scores_negative = tfidf_scores_negative.sort_values(ascending=False)

# Display the top 20 terms and their scores
print("Top 20 terms with highest TF-IDF scores in negative tweets:")
display(tfidf_scores_negative.head(20))
```

Top 20 terms with highest TF-IDF scores in negative tweets:

	0
<b>flight</b>	433.050935
<b>cancelled</b>	199.923383
<b>service</b>	169.135761
<b>hold</b>	168.683727
<b>hours</b>	168.320074
<b>customer</b>	147.458089
<b>help</b>	144.685472
<b>time</b>	136.412815
<b>delayed</b>	133.534531
<b>plane</b>	129.621976
<b>hour</b>	124.038666
<b>flightled</b>	119.962833
<b>ca</b>	117.141493
<b>bag</b>	106.595413
<b>amp</b>	106.444929
<b>gate</b>	103.755449
<b>waiting</b>	102.625705
<b>flights</b>	100.759696
<b>phone</b>	99.908555
<b>late</b>	97.032667

**dtype:** float64

```
import matplotlib.pyplot as plt
import seaborn as sns

# Extract the top 20 terms and their TF-IDF scores
top_n = 20
top_terms = tfidf_scores_negative.head(top_n)

# Create a bar plot
plt.figure(figsize=(12, 6))
sns.barplot(x=top_terms.index, y=top_terms.values, palette='viridis')

# Set labels and title
plt.xlabel('Term')
plt.ylabel('TF-IDF Score')
plt.title(f'Top {top_n} TF-IDF Terms for Negative Sentiment')

# Rotate x-axis labels for better readability
```

```
plt.xticks(rotation=45, ha='right')

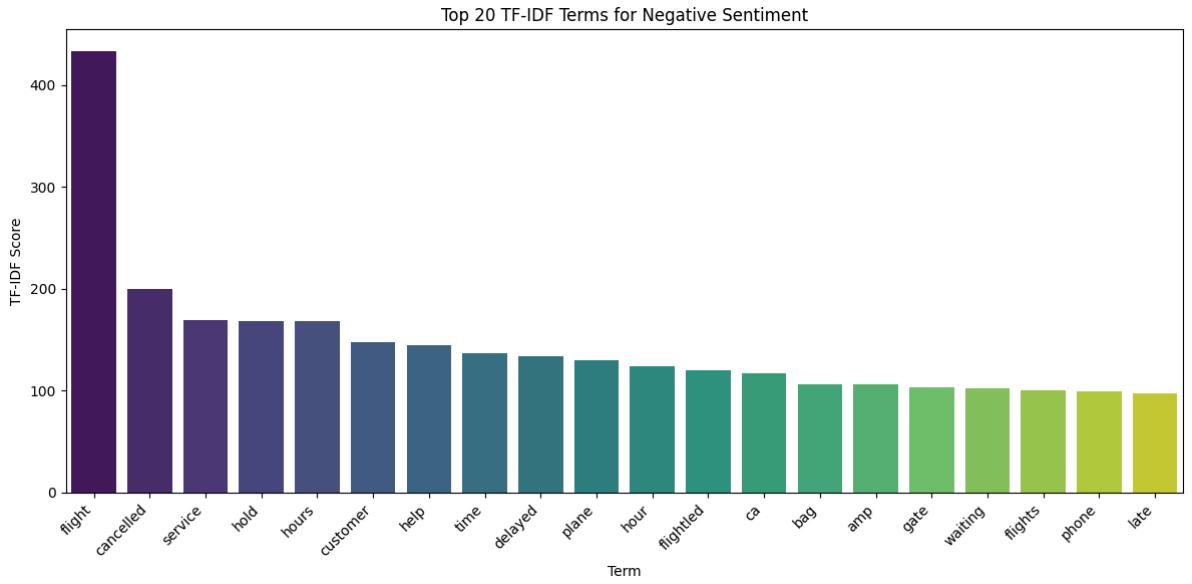
# Adjust plot parameters for a tight layout
plt.tight_layout()

# Display the plot
plt.show()

/tmp/ipython-input-1961476244.py:10: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v

```
sns.barplot(x=top_terms.index, y=top_terms.values, palette='viridis')
```



```

import matplotlib.pyplot as plt
import seaborn as sns
# Extract the top 20 terms and their TF-IDF scores
top_n = 20
top_terms = tfidf_scores_negative.head(top_n)

# Create a bar plot
plt.figure(figsize=(12, 6))
sns.barplot(x=top_terms.index, y=top_terms.values, hue=top_terms.index, palette='viridis')
# Set labels and title
plt.xlabel('Term')
plt.ylabel('TF-IDF Score')
plt.title(f'Top {top_n} TF-IDF Terms for Negative Sentiment')
# Rotate x-axis labels for better readability
plt.xticks(rotation=45, ha='right')

```

```

from wordcloud import WordCloud
import matplotlib.pyplot as plt

print("WordCloud library imported successfully.")

WordCloud library imported successfully.

```

```

wordcloud_dict = tfidf_scores_negative.to_dict()

wordcloud = WordCloud(width=800, height=400, background_color='white', colormap='viridis')

plt.figure(figsize=(15, 7))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.title('Word Cloud of Top TF-IDF Terms for Negative Sentiment')
plt.show()

```

